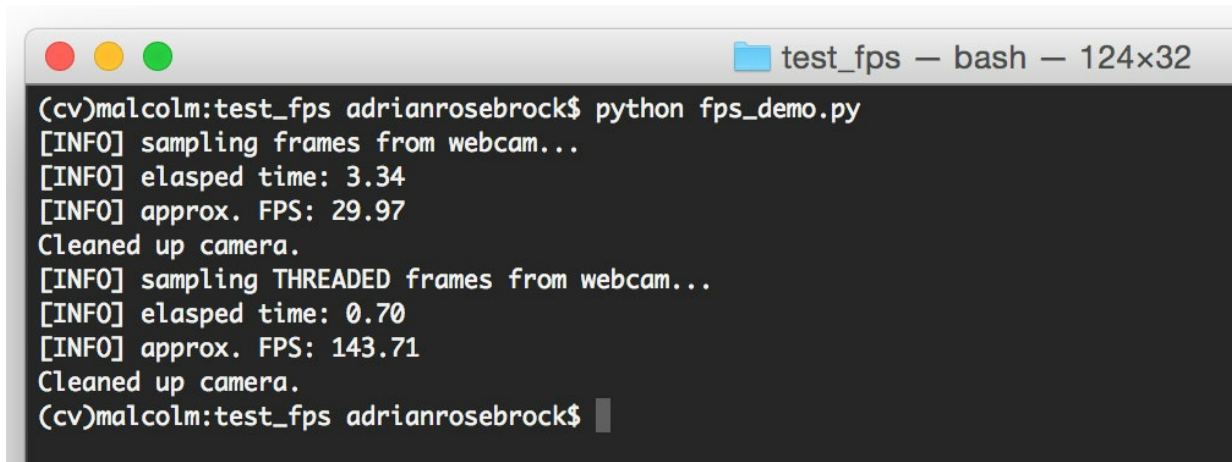


# Increasing webcam FPS with Python and OpenCV

 [pyimagesearch.com/2015/12/21/increasing-webcam-fps-with-python-and-opencv](http://pyimagesearch.com/2015/12/21/increasing-webcam-fps-with-python-and-opencv)

Adrian Rosebrock

December 21, 2015

A terminal window titled 'test\_fps — bash — 124x32' showing the execution of a Python script. The script first samples frames from a webcam, reporting an elapsed time of 3.34 seconds and an approximate FPS of 29.97. It then cleans up the camera and samples threaded frames, reporting an elapsed time of 0.70 seconds and an approximate FPS of 143.71. The terminal output is as follows:

```
(cv)malcolm:test_fps adrianrosebrock$ python fps_demo.py
[INFO] sampling frames from webcam...
[INFO] elapsed time: 3.34
[INFO] approx. FPS: 29.97
Cleaned up camera.
[INFO] sampling THREADED frames from webcam...
[INFO] elapsed time: 0.70
[INFO] approx. FPS: 143.71
Cleaned up camera.
(cv)malcolm:test_fps adrianrosebrock$
```

Over the next few weeks, I'll be doing a series of blog posts on how to ***improve your frames per second (FPS) from your webcam*** using Python, OpenCV, and threading.

Using threading to handle I/O-heavy tasks (such as reading frames from a camera sensor) is a programming model that has existed for decades.

For example, if we were to build a web crawler to spider a series of webpages (a task that is, by definition, I/O bound), our main program would spawn *multiple threads* to handle downloading the set of pages *in parallel* instead of relying on only a *single thread* (our “main thread”) to download the pages in *sequential order*. Doing this allows us to spider the webpages substantially faster.

The same notion applies to computer vision and reading frames from a camera — **we can improve our FPS simply by creating a new thread that does nothing but poll the camera for new frames while our main thread handles processing the current frame.**

This is a simple concept, but it's one that's rarely seen in OpenCV examples since it does add a few extra lines of code (or sometimes *a lot* of lines, depending on your threading library) to the project. Multithreading can also make your program harder to debug, but once you get it right, you can dramatically improve your FPS.

We'll start off this series of posts by writing a threaded Python class to access your webcam or USB camera using OpenCV.

Next week we'll use threads to improve the FPS of your Raspberry Pi and the [picamera](#) module.

Finally, we'll conclude this series of posts by creating a class that unifies *both* the threaded webcam/USB camera code and the threaded

`picamera`

code into a ***single class***, making all webcam/video processing examples on PyImageSearch *not only run faster, but run on either your laptop/desktop or the Raspberry Pi without changing a single line of code!*



Looking for the source code to this post?

[Jump Right To The Downloads Section](#) →

## Use threading to obtain higher FPS

The “secret” to obtaining higher FPS when processing video streams with OpenCV is to move the I/O (i.e., the reading of frames from the camera sensor) to a separate thread.

You see, accessing your webcam/USB camera using the

`cv2.VideoCapture`  
function and the  
`.read()`

method is a *blocking operation*. The main thread of our Python script is completely blocked (i.e., “stalled”) until the frame is read from the camera device and returned to our script.

I/O tasks, as opposed to CPU bound operations, tend to be quite slow. While computer vision and video processing applications are certainly quite CPU heavy (especially if they are intended to run in real-time), it turns out that camera I/O can be a huge bottleneck as well.

As we'll see later in this post, just by adjusting the the camera I/O process, ***we can increase our FPS by as much as 379%!***

Of course, this isn't a *true increase* of FPS as it is a *dramatic reduction in latency* (i.e., a frame is always available for processing; we don't need to poll the camera device and wait for the I/O to complete). Throughout the rest of this post, I will refer to our metrics as an “FPS increase” for brevity, but also keep in mind that it's a *combination* of both a decrease in latency and an increase in FPS.

In order to accomplish this FPS increase/latency decrease, our goal is to move the reading of frames from a webcam or USB device to an *entirely different thread, totally separate from our main Python script*.

This will allow frames to be read *continuously* from the I/O thread, all while our root thread processes the current frame. Once the root thread has finished processing its frame, it simply needs to grab the current frame from the I/O thread. This is accomplished *without* having to wait for blocking I/O operations.

The first step in implementing our threaded video stream functionality is to define a FPS

class that we can use to measure our frames per second. This class will help us obtain quantitative evidence that threading does indeed increase FPS. We'll then define a

WebcamVideoStream

class that will access our webcam or USB camera in a threaded fashion. Finally, we'll define our driver script,

fps\_demo.py

, that will compare single threaded FPS to multi-threaded FPS.

**Note:** Thanks to [Ross Milligan](#) and his blog who inspired me to do this blog post.

I've actually already implemented webcam/USB camera and

picamera

threading inside the [imutils](#) library. However, I think a discussion of the implementation can greatly improve our knowledge of *how* and *why* threading increases FPS.

To start, if you don't already have

imutils

installed, you can install it using

pip

:

Increasing webcam FPS with Python and OpenCV

\$ pip install imutils

Otherwise, you can upgrade to the latest version via:

Increasing webcam FPS with Python and OpenCV

\$ pip install --upgrade imutils

As I mentioned above, the first step is to define a

FPS

class that we can use to approximate the frames per second of a given camera + computer vision processing pipeline:

## Increasing webcam FPS with Python and OpenCV

```
# import the necessary packages
import datetime
class FPS:
    def __init__(self):
        # store the start time, end time, and total number of frames
        # that were examined between the start and end intervals
        self._start = None
        self._end = None
        self._numFrames = 0
    def start(self):
        # start the timer
        self._start = datetime.datetime.now()
        return self
    def stop(self):
        # stop the timer
        self._end = datetime.datetime.now()
    def update(self):
        # increment the total number of frames examined during the
        # start and end intervals
        self._numFrames += 1
    def elapsed(self):
        # return the total number of seconds between the start and
        # end interval
        return (self._end - self._start).total_seconds()
    def fps(self):
        # compute the (approximate) frames per second
        return self._numFrames / self.elapsed()
```

On **Line 5-10** we define the constructor to our

FPS

class. We don't require any arguments, but we do initialize three important variables:

- `_start`  
: The starting timestamp of when we commenced measuring the frame read.
- `_end`  
: The ending timestamp of when we stopped measuring the frame read.
- `_numFrames`  
: The total number of frames that were read during the  
`_start`  
and  
`_end`  
interval.

**Lines 12-15** define the

start

method, which as the name suggests, kicks-off the timer.

Similarly, **Lines 17-19** define the

stop

method which grabs the ending timestamp.

The

update

method on **Lines 21-24** simply increments the number of frames that have been read during the starting and ending interval.

We can grab the total number of seconds that have elapsed between the starting and ending interval on **Lines 26-29** by using the

elapsed

method.

And finally, we can approximate the FPS of our camera + computer vision pipeline by using the

fps

method on **Lines 31-33**. By taking the total number of frames read during the interval and dividing by the number of elapsed seconds, we can obtain our estimated FPS.

Now that we have our

FPS

class defined (so we can empirically compare results), let's define the WebcamVideoStream

class which encompasses the actual threaded camera read:

Increasing webcam FPS with Python and OpenCV

```
# import the necessary packages
```

```
from threading import Thread
```

```
import cv2
```

```
class WebcamVideoStream:
```

```
def __init__(self, src=0):
```

```
# initialize the video camera stream and read the first frame  
# from the stream
```

```
self.stream = cv2.VideoCapture(src)
```

```
(self.grabbed, self.frame) = self.stream.read()
```

```
# initialize the variable used to indicate if the thread should  
# be stopped
```

```
self.stopped = False
```

We define the constructor to our

WebcamVideoStream

class on **Line 6**, passing in an (optional) argument: the

src  
of the stream.  
If the

src  
is an *integer*, then it is presumed to be the *index* of the webcam/USB camera on your system. For example, a value of  
src=0  
indicates the *first* camera and a value of  
src=1  
indicates the *second* camera hooked up to your system (provided you have a second one, of course).  
If

src  
is a *string*, then it assumed to be the *path* to a video file (such as .mp4 or .avi) residing on disk.  
**Line 9** takes our

src  
value and makes a call to  
cv2.VideoCapture  
which returns a pointer to the camera/video file.  
Now that we have our

stream  
pointer, we can call the  
.read()  
method to poll the stream and grab the next available frame ( **Line 10**). This is done strictly for initialization purposes so that we have an *initial* frame stored in the class. We'll also initialize

stopped  
, a boolean indicating whether the threaded frame reading should be stopped or not. Now, let's move on to actually utilizing threads to read frames from our video stream using OpenCV:

Increasing webcam FPS with Python and OpenCV  
def start(self):  
# start the thread to read frames from the video stream  
Thread(target=self.update, args=()).start()  
return self  
def update(self):  
# keep looping infinitely until the thread is stopped  
while True:  
# if the thread indicator variable is set, stop the thread

```

if self.stopped:
    return
# otherwise, read the next frame from the stream
(self.grabbed, self.frame) = self.stream.read()
def read(self):
    # return the frame most recently read
    return self.frame
def stop(self):
    # indicate that the thread should be stopped
    self.stopped = True
Lines 16-19 define our

```

**start**  
 method, which as the name suggests, starts the thread to read frames from our video stream. We accomplish this by constructing a **Thread** object using the **update** method as the callable object invoked by the **run()** method of the thread.  
 Once our driver script calls the

**start**  
 method of the **WebcamVideoStream** class, the **update** method (**Lines 21-29**) will be called.  
 As you can see from the code above, we start an infinite loop on **Line 23** that continuously reads the next available frame from the video

stream  
 via the **.read()** method (**Line 29**). If the **stopped** indicator variable is ever set, we break from the infinite loop (**Lines 25 and 26**).  
 Again, keep in mind that once the

**start**  
 method has been called, the **update** method is placed in a *separate thread from our main Python script* — **this separate thread is how we obtain our increased FPS performance.**

In order to access the most recently polled

frame

from the

stream

, we'll use the

read

method on **Lines 31-33**.

Finally, the

stop

method (**Lines 35-37**) simply sets the  
stopped

indicator variable and signifies that the thread should be terminated.

Now that we have defined both our

FPS

and

WebcamVideoStream

classes, we can put all the pieces together inside  
fps\_demo.py

:

Increasing webcam FPS with Python and OpenCV

# import the necessary packages

from \_\_future\_\_ import print\_function

from imutils.video import WebcamVideoStream

from imutils.video import FPS

import argparse

import imutils

import cv2

# construct the argument parse and parse the arguments

ap = argparse.ArgumentParser()

ap.add\_argument("-n", "--num-frames", type=int, default=100,  
help="# of frames to loop over for FPS test")

ap.add\_argument("-d", "--display", type=int, default=-1,  
help="Whether or not frames should be displayed")

args = vars(ap.parse\_args())

We start off by importing our necessary packages on **Lines 2-7**. Notice how we are  
importing the

FPS

and

WebcamVideoStream

classes from the imutils library. If you do not have  
imutils



installed or you need to upgrade to the latest version, please see the note at the top of this section.

**Lines 10-15** handle parsing our command line arguments. We'll require two switches here:

--num-frames

, which is the number of frames to loop over to obtain our FPS estimate, and

--display

, an indicator variable used to specify if we should use the

cv2.imshow

function to display the frames to our monitor or not.

The

--display

argument is actually really important when approximating the FPS of your video processing pipeline. Just like reading frames from a video stream is a form of I/O, *so is displaying the frame to your monitor!* We'll discuss this in more detail inside the **Threading results** section of this post.

Let's move on to the next code block which does *no threading* and uses *blocking I/O* when reading frames from the camera stream. This block of code will help us obtain a *baseline* for our FPS:

Increasing webcam FPS with Python and OpenCV

```
# grab a pointer to the video stream and initialize the FPS counter
```

```
print("[INFO] sampling frames from webcam...")
```

```
stream = cv2.VideoCapture(0)
```

```
fps = FPS().start()
```

```
# loop over some frames
```

```
while fps._numFrames < args["num_frames"]:
```

```
# grab the frame from the stream and resize it to have a maximum
```

```
# width of 400 pixels
```

```
(grabbed, frame) = stream.read()
```

```
frame = imutils.resize(frame, width=400)
```

```
# check to see if the frame should be displayed to our screen
```

```
if args["display"] > 0:
```

```
cv2.imshow("Frame", frame)
```

```
key = cv2.waitKey(1) & 0xFF
```

```
# update the FPS counter
```

```
fps.update()
```

```
# stop the timer and display FPS information
```

```
fps.stop()
```

```
print("[INFO] elapsed time: {:.2f}".format(fps.elapsed()))
```

```
print("[INFO] approx. FPS: {:.2f}".format(fps.fps()))
```

```
# do a bit of cleanup
```

```
stream.release()
```

```
cv2.destroyAllWindows()
```

**Lines 19 and 20** grab a pointer to our video stream and then start the FPS counter.

We then loop over the number of desired frames on **Line 23**, read the frame from camera (**Line 26**), update our FPS counter (**Line 35**), and optionally display the frame to our monitor (**Lines 30-32**).

After we have read

```
--num-frames
```

from the stream, we stop the FPS counter and display the elapsed time along with approximate FPS on **Lines 38-40**.

Now, let's look at our *threaded* code to read frames from our video stream:

Increasing webcam FPS with Python and OpenCV

```
# created a *threaded* video stream, allow the camera sensor to warmup,
```

```
# and start the FPS counter
```

```
print("[INFO] sampling THREADED frames from webcam...")
```

```
vs = WebcamVideoStream(src=0).start()
```

```
fps = FPS().start()
```

```
# loop over some frames...this time using the threaded stream
```

```
while fps._numFrames < args["num_frames"]:
```

```
# grab the frame from the threaded video stream and resize it
```

```
# to have a maximum width of 400 pixels
```

```
frame = vs.read()
```

```
frame = imutils.resize(frame, width=400)
```

```
# check to see if the frame should be displayed to our screen
```

```
if args["display"] > 0:
```

```
cv2.imshow("Frame", frame)
```

```
key = cv2.waitKey(1) & 0xFF
```

```
# update the FPS counter
```

```
fps.update()
```

```
# stop the timer and display FPS information
```

```
fps.stop()
```

```
print("[INFO] elapsed time: {:.2f}".format(fps.elapsed()))
```

```
print("[INFO] approx. FPS: {:.2f}".format(fps.fps()))
```

```
# do a bit of cleanup
```

```
cv2.destroyAllWindows()
```

```
vs.stop()
```

Overall, this code looks near identical to the code block above, only this time we are leveraging the

```
WebcamVideoStream
```

```
class.
```

We start the threaded stream on **Line 49**, loop over the desired number of frames on **Lines 53-65** (again, keeping track of the total number of frames read), and then display our output on **Lines 69 and 70**.

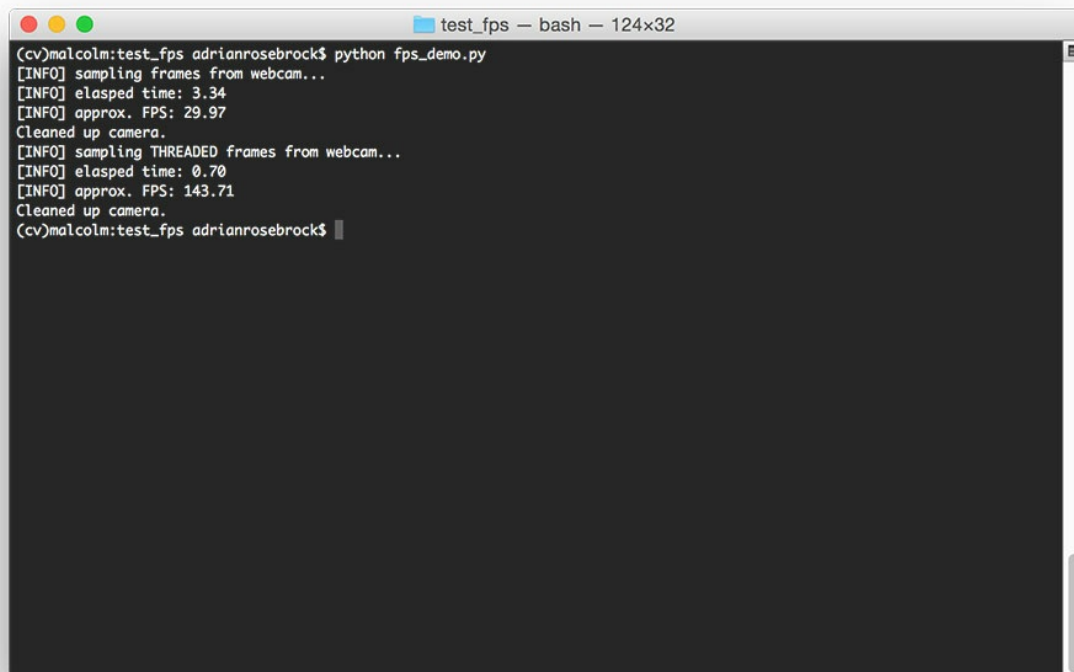
## Threading results

---

To see the affects of webcam I/O threading in action, just execute the following command:

Increasing webcam FPS with Python and OpenCV

```
$ python fps_demo.py
```



```
(cv)malcolm:test_fps adrianrosebrock$ python fps_demo.py
[INFO] sampling frames from webcam...
[INFO] elapsed time: 3.34
[INFO] approx. FPS: 29.97
Cleaned up camera.
[INFO] sampling THREADED frames from webcam...
[INFO] elapsed time: 0.70
[INFO] approx. FPS: 143.71
Cleaned up camera.
(cv)malcolm:test_fps adrianrosebrock$
```

**Figure 1:** By using threading with Python and OpenCV, ***we are able to increase our FPS by over 379%!***

As we can see, by using no threading and sequentially reading frames from our video stream in the *main thread* of our Python script, we are able to obtain a respectable 29.97 FPS.

However, once we switch over to using *threaded camera I/O*, we reach 143.71 FPS — ***an increase of over 379%!***

This is clearly a *huge decrease* in our latency and a *dramatic increase* in our FPS, obtained simply by using threading.

However, as we're about to find out, using the

```
cv2.imshow
```

can *substantially decrease* our FPS. This behavior makes sense if you think about it — the

`cv2.show`

function is just another form of I/O, only this time instead of reading a frame from a video stream, we're instead sending the frame to output on our display.

**Note:** We're also using the

`cv2.waitKey(1)`

function here which does add a 1ms delay to our main loop. That said, this function is necessary for keyboard interaction and to display the frame to our screen (especially once we get to the Raspberry Pi threading lessons).

To demonstrate how the

`cv2.imshow`

I/O can decrease FPS, just issue this command:

Increasing webcam FPS with Python and OpenCV

```
$ python fps_demo.py --display 1
```

**Figure 2:** Using the `cv2.imshow` function can reduce our FPS — it is another form of I/O, after all!

Using *no threading*, we reach 28.90 FPS. And *with threading* we hit 39.93 FPS. This is still a **38% increase** in FPS, but nowhere near the **379% increase** from our previous example.

Overall, I recommend using the

`cv2.imshow`

function to help debug your program — but if your final production code doesn't need it, *there is no reason to include it* since you'll be hurting your FPS.

A great example of such a program would be developing a home surveillance motion detector that sends you a txt message containing a photo of the person who just walked in the front door of your home. Realistically, you do not need the

`cv2.imshow`

function for this. By removing it, you can increase the performance of your motion detector and allow it to process more frames faster.

## Summary

---


In this blog post we learned how threading can be used to increase your webcam and USB camera FPS using Python and OpenCV.

As the examples in this post demonstrated, we were able to obtain a **379% increase in FPS** simply by using threading. While this isn't necessarily a fair comparison (since we could be processing the same frame multiple times), it does demonstrate the importance of *reducing latency* and always having a frame ready for processing.

**In nearly all situations, using threaded access to your webcam can substantially improve your video processing pipeline.**

Next week we'll learn how to increase the FPS of our Raspberry Pi using the [picamera](#) module.

***Be sure to enter your email address in the form below to be notified when the next post goes live!***



## Download the Source Code and FREE 17-page Resource Guide

Enter your email address below to get a .zip of the code and a **FREE 17-page Resource Guide on Computer Vision, OpenCV, and Deep Learning**. Inside you'll find my hand-picked tutorials, books, courses, and libraries to help you master CV and DL!